



1 Problem definition (1–2 minutes)

- Separate backend services per Salesforce object (e.g., Lead, Opportunity, Account), and each service integrated with Salesforce differently
- Integration was tight and synchronous, so failures cascaded and debugging was painful
- Burned engineering time without improving quality or user value
- Maintenance cost was high because there was no common platform or shared integration pattern across services
- Training implemented differently per object, which led to inconsistent user experience and uneven model quality

Goal

Build a single, stable integration surface for Salesforce and move heavy compute to an async, tenant-isolated ML platform with reliable training + prediction and clear model lifecycle.

2 Functional requirements (3–4 minutes)

- Tenant onboarding (create tenant namespace + metadata)
- Consent management
 - allowTraining (use data for training)
 - allowStorePredictions (persist outcomes)
- Bulk export from Salesforce (per tenant) into staging (S3/lakehouse)
- Training orchestration per tenant (start, retry, resume, cancel)
- Store artifacts + metrics, register model versions, promote active model
- On-demand prediction for a Salesforce record (lead/opportunity/account etc.)
- Support both sync and async prediction modes
- Write results back into Salesforce objects/fields and/or Salesforce-hosted store
- Observability and admin APIs for status, run history, and failures

3 Non-functional requirements (2–3 minutes)

Highlight only what drives architecture:

- High availability for prediction path (99.9% target)
- Eventual consistency acceptable for training and async prediction callbacks
- Multi-tenant isolation (data + model + access)
- Burst handling (spiky predictions, bulk exports)
- Idempotency and resumability (exports, training runs, prediction requests)
- Compliance: consent gating must be enforced server-side
- Encryption in transit + at rest, least privilege IAM

This justifies **single integration layer + event-driven queues + orchestrated ML pipeline**.

4 APIs (2 minutes)

Keep it short and interview-ready:

Tenant + consent

- POST /v1/tenants
- POST /v1/tenants/{tenantId}/consent
- GET /v1/tenants/{tenantId}/status

Export + training

- POST /v1/tenants/{tenantId}/exports
- GET /v1/tenants/{tenantId}/exports/{exportJobId}
- POST /v1/tenants/{tenantId}/training-runs
- GET /v1/tenants/{tenantId}/training-runs/{trainingRunId}

Prediction

- Sync: POST /v1/tenants/{tenantId}/predict
- Async: POST /v1/tenants/{tenantId}/predict:async +
GET /v1/predictions/{requestId}

Callbacks to Salesforce (internal integration)

- POST /v1/salesforce/callbacks/prediction-complete
- POST /v1/salesforce/callbacks/model-ready

Mention **direct multi-service SF integration only to reject it**. Shows judgment.

5 Data model + partitioning (5 minutes)

Explain the minimum entities that make the system real:

- **Tenant**: tenant_id , sf_org_id , status , timestamps
- **Consent**: tenant_id , allow_training , allow_store_predictions , updated_by , version
- **ExportJob**: export_job_id , tenant_id , sf_bulk_job_id , dataset_uri , schema_hash , status
- **TrainingRun**: training_run_id , tenant_id , export_job_id , status , sagemaker_job_arn , metrics_uri , model_artifact_uri
- **ModelRegistry**: tenant_id , model_version , status (staged/active/deprecated) , artifact_uri
- **PredictionRequest**: request_id (idempotency key), tenant_id , entity_type , entity_id , status , model_version_used , result_uri

Partitioning / isolation (this is the key point):

- **Primary partition key**: tenant_id
- Storage layout: s3://.../{tenant_id}/{export_job_id}/... and .../{model_version}/...
- DB indexes: (tenant_id, status) , (tenant_id, created_at) , (tenant_id, entity_id)

This naturally leads into HLD.

6 High-level design (5–7 minutes)

You already have the right architecture. Walk through **one clean training path** and **one clean prediction path**.

Training path (write-heavy, async)

- Salesforce (Einstein/app) → API Gateway
- Tenant Service validates consent + config
- Export Worker uses SF Bulk API → writes dataset to S3/lakehouse
- Orchestrator creates TrainingRun → sends message to `SQS.training`
- Workers/Lambda start SageMaker training job
- Artifacts + metrics → S3
- ModelRegistry updated → new active model (optional gating)
- Callback worker updates Salesforce (model-ready + metadata)

Prediction path (read/compute, sync + async)

- Salesforce → API Gateway → Inference Service
- Resolve tenant active model version
- **Sync:** call warm model endpoint or cached runtime → return prediction
- **Async:** enqueue `SQS.inference` → worker runs prediction → callback writes back to Salesforce

Don't enumerate every box. Focus on flow and why boundaries exist.

7 Deep dive (10–15 minutes)

Slow down here. Pick **1–2 problems** that show engineering depth.

Deep dive A: Consent-gated training correctness under retries (the “hard” one)

Cover only:

- **Server-side consent gate**
 - Orchestrator checks `Consent.allow_training == true` at training-run creation

time

- Option: re-check right before starting SageMaker job (protect against mid-flight revoke)

- **Idempotency**

- `trainingRunId` is the idempotency key
- Workers must be able to receive duplicates and do nothing if already `Succeeded/Running`

- **At-least-once delivery**

- SQS may deliver duplicate messages
- Worker logic: read TrainingRun state → only start job if state == `Queued`

- **Resumability**

- ExportJob and TrainingRun statuses are persisted
- If export fails, retry export without duplicating datasets (use `exportJobId` path)

- **DLQ + replay**

- Poison messages go to DLQ with reason
- Operator can replay after fixing config/schema mismatch

This is where interviewers engage because it's real failure-mode thinking.

Deep dive B: Prediction spike control and callback reliability

Cover only:

- **Sync vs async switching**

- Use API Gateway throttles per tenant
- If queue depth high or endpoint saturated → force async

- **Model loading strategy**

- Warm endpoints for top tenants (latency)
- Lambda runtime load for long tail (cost), accept cold-start
- Hybrid: keep N hottest models warm

- **Request idempotency**

- `requestId` required from Salesforce
- if same requestId repeats → return stored result or current status

- **Callback delivery**

- Salesforce write-back can fail (rate limits, transient errors)
- Use retry with backoff + “pending callback” status

- Persist last error + next retry time for ops visibility
-

8 Tradeoffs and closing (2 minutes)

End strong:

- Async pipeline adds backend complexity (orchestrator, queues, retries)
- Eventual consistency for training and async prediction callbacks
- Some latency/cost tradeoffs depending on warm endpoints vs Lambda model loads
- But you get:
 - Cleaner Salesforce integration (single surface)
 - Much better reliability under spikes
 - Tenant isolation + consent enforcement
 - Real model lifecycle and auditability
 - Easier evolution over time (new models/features without breaking SF)